

# Python tools for music analysis

Brian McFee



# Why Python?

- Sooner or later, we're all software engineers...
  - ... whether we want to be or not
- Python offers a nice balance between power and simplicity
- Active developer community
- Also it's free!

# Why not python?

- 20+ years of entrenched legacy code and tools
- Change is hard! Porting code is hard!
  - *Just one more project in matlab...*
- Relatively little support for audio-related tasks
  - *... but I'm trying to change that!*

# The (scientific) python ecosystem

(just the highlight reel)

# The core: numpy and scipy

- **numpy**
  - numerical array data structures
  - basic mathematical operations
- **scipy**
  - linear algebra, sparse matrices
  - optimization, DSP, statistics

```
>>> import numpy as np
>>> x = np.arange(5)
>>> x
array([0, 1, 2, 3, 4])
>>> x**2
array([0, 1, 4, 9, 16])
```

```
>>> import scipy
>>> scipy.fft(x)
array([ 10.0+0.j , -2.5+3.441j,
       -2.5+0.812j, -2.5-0.812j,
       -2.5-3.441j])
```

[\[numpy for matlab users\]](#)

# Visualization and analysis

- **matplotlib**
  - powerful, but difficult to use properly
- **seaborn**
  - makes matplotlib nice by default
  - adds statistical functionality
- **pandas**
  - time series and statistics (Series)
  - named fields (DataFrames)
  - relational structures (pivots, joins, etc)

```
>>> import matplotlib.pyplot as plt
>>> fig = plt.figure()
>>> plt.plot(x, x**2)
>>> fig.show()
```

```
>>> import seaborn as sns
>>> sns.set()
>>> # [Repeat the plot code above]
```

```
>>> import pandas as pd
>>> data = pd.read_csv('file.csv')
>>> data.describe()
```

# IPython and IPython notebook

- Interactive python shell
  - Always use `ipython` instead of `python` shell

- IPython Notebook
  - Code *interactively the browser!*
    - images (and plots)
    - video and audio playback
    - widgets and controls
  - Share notebooks on [nbviewer](#)

- [Jupyter](#) = Notebook - Python

```
In [4]: import librosa

In [12]: y, sr = librosa.load('data/SALAMI/audio/1216.mp3')

In [13]: tempo, beats = librosa.beat.beat_track(y, sr=sr)

In [14]: print('Tempo: {:.2f} BPM'.format(tempo))
          print('First 10 beats (s): \n {:.s}'.format(librosa.frames_to_time(beats[
```

Tempo: 112.35 BPM  
First 10 beats (s):  
[ 0.209 0.743 1.277 1.834 2.368 2.926 3.46 3.994 4.528 5.085]

# The (musical) python ecosystem



# Symbolic data

- **music21**
  - process symbolic score data
  - MusicXML and humdrum
  - [extensive documentation](#)
- **pretty\_midi**
  - simplified access to midi data
  - easy to read and write
  - [documentation / examples](#)

```
>>> from music21 import corpus  
>>> sBach = corpus.parse('bach/bwv7.7')  
>>> sBach.show()
```

# Music information retrieval

- **mir\_eval**
  - *How good is my algorithm?*
  - standardized implementation of evaluation metrics
  - stable release, but looking to expand
- **JAMS**
  - JSON storage for annotations
  - collect all annotations for a track in one place
  - standardized schema, validation
  - under (very) active development

# librosa: audio processing in python

- Tools to build your own MIR system
  - decode audio
  - STFT, CQT, Mel, chroma, ...
  - source decomposition, HPSS
  - onset detection, beat tracking, structural features, ...
  - spectrogram visualization
- Provides features for various tasks:
  - structure analysis, similarity, chord and instrument rec, etc
- Some integration with **scikit-learn**

# Enough talk, let's code!

Stable version:

```
$ pip install --user librosa
```

Development version (recommended):

<https://github.com/bmcfee/librosa/releases>

# Notebook links

- [librosa demo](#)
- [Widgets](#)
- [Loudness vs IOI](#)

# But how do I make it go fast?

- Loops can be slow, just like in matlab
- First step: profile your code!
  - `$ python -m cProfile -o your_program.profile your_program.py`
  - `$ runsnake your_program.profile`
- Solution 1: vectorize your loops!
- Solution 2: **numba** [automagic]
- Solution 3: **cython** [semi-automagic]
- Solution 4: **scipy.weave** [embed C/C++]
- Solution 5: **joblib.Parallel** [use multiple threads]

# What are the drawbacks?

- Real-time audio is still a ways out
  - [pyaudio](#) gives PortAudio bindings
  - latency is still pretty high, but **cython** might help
- Python 2.7 vs 3
  - If you can, just use **python3**
  - **six** can make the transition easier
- Development is rapid, and it can be hard to keep up

# Parting words / best practices

- Use **pylint** and **pep8** to avoid simple mistakes
- Use **nose** to handle testing
- Use the **six** module for py2/py3 compatibility
  - even if you only use one version, play nice with others!
- Use **sphinx** and **numpydoc** for documentation
- Use **notebook** for reproducibility, and version control everything!



# Thanks!

# Questions?

[brian.mcfree@nyu.edu](mailto:brian.mcfree@nyu.edu)  
<http://bmcfree.github.io/>

# References / packages

- [numpy](#)
- [scipy](#)
- [matplotlib](#)
- [seaborn](#)
- [bokeh](#)
- [mpld3](#)
- [pandas](#)
- [anaconda](#)
- [ipython](#)
- [scikit-learn](#)
- [music21](#)
- [pretty\\_midi](#)
- [mir\\_eval](#)
- [jams](#)
- [librosa](#)
- [pysoundfile](#)
- [scikits.samplerate](#)
- [cProfile](#)
- [runsnake](#)
- [numba](#)
- [cython](#)
- [scipy.weave](#)
- [joblib.Parallel](#)
- [six](#)
- [pylint](#)
- [pep8](#)
- [nose](#)
- [sphinx](#)
- [numpydoc](#)